

Oct 02, 18 19:57	cifar10.py	Page 1/2
<pre>''' ECE471, Selected Topics in Machine Learning – Assignment 4 Submit by Oct. 4, 10pm tldr: Classify cifar10. Achieve performance similar to the state of the art. Classify cifar100. Achieve a top-5 accuracy of 70%. The architecture is adapted from the AllConv net presented in this paper: https://arxiv.org/pdf/1412.6806.pdf And some portions of code were modified from this GitHub repo: https://github.com/MateLabs/All-Conv-Keras My contributions: - Experimented with placement and probabilities of dropout layers (removed the first dropout layer, contrary to paper). - Experimented with optimizers (decided that SGD was still the best, but used Nesterov momentum and a different learning rate, contrary to paper). - Trained it for longer, contrary to GitHub repo. ''' import numpy as np import tensorflow as tf from tensorflow import keras from tensorflow.keras.datasets.cifar10 import load_data from tensorflow.keras.preprocessing.image import ImageDataGenerator # Data parameters NUM_CLASSES = 10 HEIGHT = 32 WIDTH = 32 NUM_CHANNELS = 3 # Load data and split into train, val, test sets (x_train, y_train), (x_test, y_test) = load_data() (x_val, y_val), (x_test, y_test) = \ (x_test[:5000], y_test[:5000]), (x_test[5000:], y_test[5000:]) # Normalize and reshape data and labels x_train, x_val, x_test = \ map(lambda x: (x / 255.0).reshape([-1, HEIGHT, WIDTH, NUM_CHANNELS]), [x_train, x_val, x_test]) y_train, y_val, y_test = \ map(lambda y: keras.utils.to_categorical(y, NUM_CLASSES), [y_train, y_val, y_test]) datagen = ImageDataGenerator(rotation_range=0.1, width_shift_range=0.1, height_shift_range=0.1, horizontal_flip=True) datagen.fit(x_train) # Hyperparameters BATCH_SIZE = 128 NUM_EPOCHS = 200 # I feel bad about this... model = keras.Sequential() # Idk if I can wrap this architecture with a function (DRY) without making # it less readable... Notice the exceptions to the patterns. model.add(keras.layers.Conv2D(96, 3, activation='relu', padding='same', input_shape=[32, 32, 3])) model.add(keras.layers.Conv2D(96, 3, activation='relu', padding='same')) model.add(keras.layers.Conv2D(96, 3, activation='relu', padding='same', strides=2)) model.add(keras.layers.Dropout(0.5)) model.add(keras.layers.Conv2D(192, 3, activation='relu', padding='same')) model.add(keras.layers.Conv2D(192, 3, activation='relu', padding='same')) model.add(keras.layers.Conv2D(192, 3, activation='relu', padding='same', strides=2))</pre>		

Oct 02, 18 19:57	cifar10.py	Page 2/2
<pre>model.add(keras.layers.Dropout(0.5)) model.add(keras.layers.Conv2D(192, 3, activation='relu', padding='same')) model.add(keras.layers.Conv2D(192, 1, activation='relu', padding='valid')) model.add(keras.layers.Conv2D(10, 1, activation='relu', padding='valid')) model.add(keras.layers.GlobalAveragePooling2D()) model.add(keras.layers.Dense(NUM_CLASSES, activation=keras.activations.softmax)) model.compile(loss=keras.losses.categorical_crossentropy, optimizer=keras.optimizers.SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True), metrics=['accuracy']) model.fit_generator(datagen.flow(x_train, y_train, batch_size=BATCH_SIZE), epochs=NUM_EPOCHS, verbose=1, validation_data=(x_val, y_val)) loss, acc = model.evaluate(x_test, y_test, verbose=1) print('Test loss:', loss) print('Test accuracy:', acc) ''' Output, omitting Keras logs. Test loss: 0.3725 Test accuracy: 0.9054 '''</pre>		